

# Query Optimization Strategy for Building Management System

## 1. Introduction

Query optimization is a crucial aspect of database management, ensuring that data retrieval is **efficient, scalable, and cost-effective**. In a **building management system (BMS)**, where multiple databases (MongoDB, SQL, Time-Series, and Vector DB) interact, an optimized query execution strategy is essential for:

- **Reducing latency** in retrieving equipment failure logs.
- **Minimizing computational overhead** for real-time anomaly detection.
- **Ensuring smooth integration** between the **query parser, executor, and databases**.

This document outlines an **effective query optimization strategy** for handling structured queries generated by an **NLP query parser** and optimizing execution across different database types.

## 2. Understanding the Query Flow in the System

The system architecture consists of the following components:

1. **NLP Query Parser (Python + OpenAI API):** Converts user queries into structured database queries.
2. **Query Executor (Python):** Processes structured queries and determines the appropriate database.
3. **Multiple Databases:**
  - **MongoDB (Cosmos DB):** Stores logs and maintenance records.
  - **SQL Database:** Maintains financial and report-related data.
  - **Time-Series Database:** Handles sensor data such as temperature, power consumption, and HVAC performance.
  - **Vector Database:** Stores document-based reports for retrieval and analysis.

## 3. Query Optimization Techniques

### 3.1. Indexing for Faster Query Execution

Indexing improves **search efficiency** by reducing the time required to fetch data from large datasets.

- **MongoDB:** Create **compound indexes** on frequently queried fields like equipment and timestamp.
- **SQL Database:** Use **B-Trees and Hash Indexes** for structured data storage.
- **Time-Series Database:** Implement **time-based partitioning** to improve retrieval speed.

- **Vector Database:** Use **approximate nearest neighbors (ANN)** for fast similarity searches.

### 3.2. Query Caching for Reducing Redundant Computation

To avoid repeated execution of the same query, **caching mechanisms** can be implemented:

- **Redis Cache:** Store frequently queried data (e.g., past equipment failures) to reduce database load.
- **Query Result Caching:** Store structured query results for a fixed duration to prevent repetitive computations.
- **Materialized Views in SQL:** Precompute frequently used aggregations (e.g., average repair time per equipment type).

### 3.3. Query Optimization Techniques Per Database Type

#### MongoDB (Logs & Maintenance)

- Use **aggregation pipelines** to **filter, group, and transform** maintenance logs efficiently.
- **Sharding strategy:** Distribute data across multiple nodes for faster retrieval.
- Implement **TTL Indexes** to automatically delete old logs beyond a retention period.

#### SQL Database (Financials & Reports)

- Use **EXPLAIN ANALYZE** to check query execution plans.
- Optimize **JOIN operations** by using indexed columns.
- Apply **denormalization** where necessary to reduce complex multi-table queries.

#### Time-Series Database (Sensor Data)

- Store data in **optimized time-based partitions** (e.g., hourly, daily).
- Use **downsampling techniques** to reduce data volume for older records.
- Implement **window functions** for fast time-based aggregations.

#### Vector Database (Document Search)

- Use **Approximate Nearest Neighbour's (ANN) indexing** for efficient document retrieval.
- Optimize **embedding similarity search** by reducing vector dimensionality.

## 4. Query Execution Plan for Different Scenarios

### Scenario 1: Retrieve HVAC Failures in the Last 3 Months

- **Optimized Approach:**
  - Use an **index scan** on equipment and timestamp fields in MongoDB.

- Apply **aggregation pipeline** to filter records within the date range.

### Scenario 2: Analyse Energy Consumption Trends Over the Last 6 Months

- **Optimized Approach:**
  - Query the **Time-Series Database** with a **time-partitioned scan**.
  - Use **rolling aggregations** to reduce computational load.
  - Implement **caching for recent queries** to improve response time.

### Scenario 3: Fetch Equipment Maintenance Costs for Q1 Reports

- **Optimized Approach:**
  - Use **indexed joins** in the SQL database to fetch financial data.
  - Precompute and store **aggregated maintenance costs** using materialized views.
  - Apply **query caching** for frequently accessed reports.

## 5. Conclusion

A **well-optimized query strategy** significantly improves the performance of the building management system by:

- **Reducing query execution time** through indexing and caching.
- **Minimizing database load** with optimized storage techniques.
- **Ensuring fast and scalable retrieval** of equipment logs, financial data, and sensor readings.

By implementing these optimization techniques, the system can handle **real-time anomaly detection** efficiently, ensuring **timely maintenance interventions** and **improving overall building management operations**.